

Parallel Computing with MATLAB

Advanced Big Data Analytics

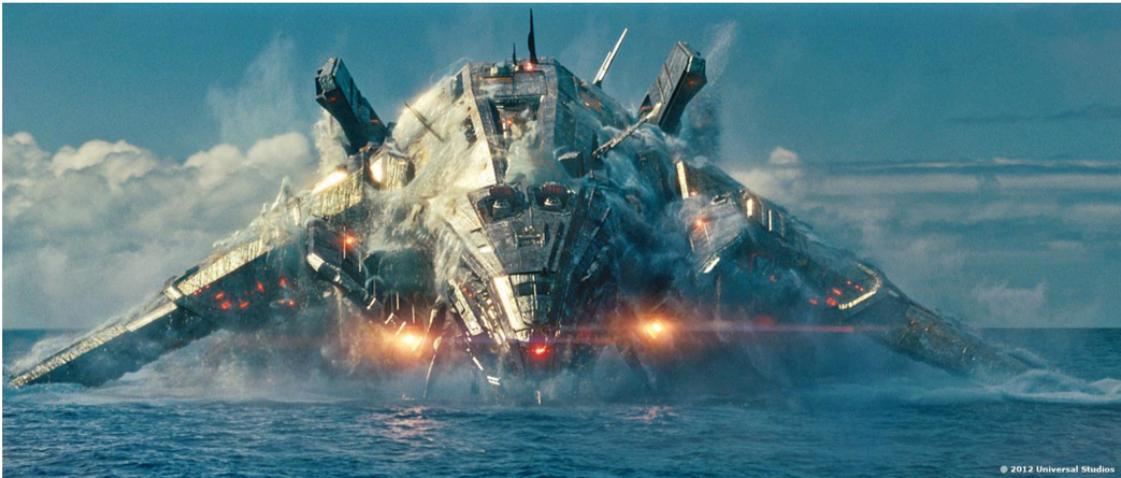
GPU in Graphics

Hollywood Visual Effects

One often cannot film various real-world situations required in order to tell a story

- Some times the situation is too dangerous, impractical, expensive, or rare
- Other times the situation doesn't actually exist, except in an alternative reality

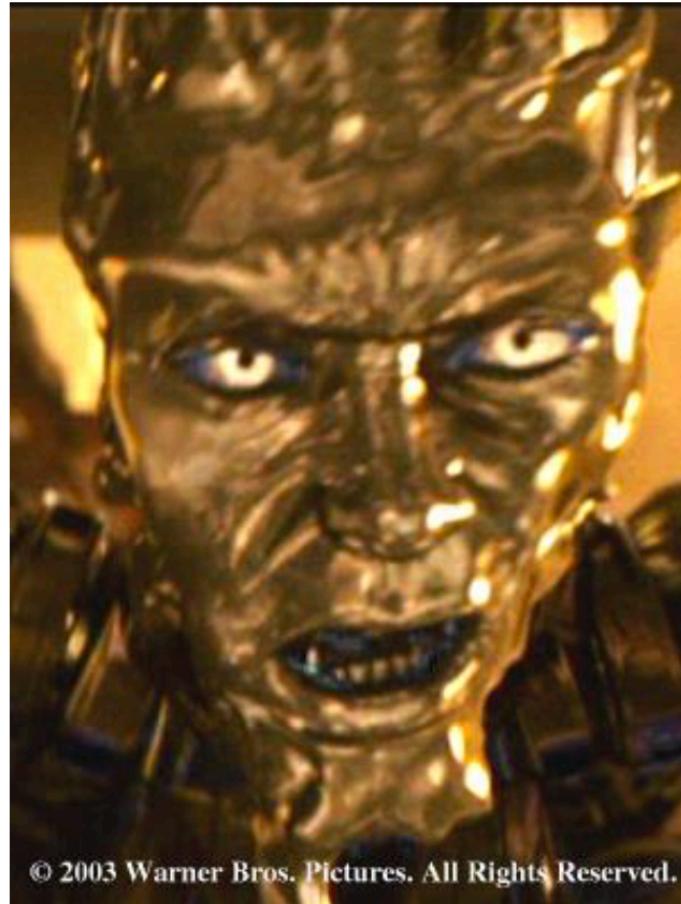
Visual Effects: Liquids



Battleship



The Day After Tomorrow



© 2003 Warner Bros. Pictures. All Rights Reserved.

Terminator 2

Visual Effects: Gases



**Harry Potter and the Order of
the Phoenix**



Terminator 3



Star Wars Episode III

SIGGRAPH 2015

<http://kesen.realtimerendering.com/sig2015.html>

SIGGRAPH 2016

<http://kesen.realtimerendering.com/sig2016.html>

Siggraph Asia

Siggraph

IEEE IEEE Transactions on

Visualization and Computer Graphics(TCVG)

Chinagraph

Visual Effects: Solids

Destruction:
fracture,
explosions,
etc.



Super 8



2012

Visual Effects: CG Creatures



Yoda, Star Wars Episode II



**Sméagol/Gollum, The Lord of
the Rings**

Motion Capture



Facial capture in Avatar



**Motion capture of Olympic
swimmer Dana Vollmer by
Manhattan Mocap
(technology transition)**

Animated Films



Toy Story 3



Monsters, Inc.

Video Games



Spore



Crysis



Braid

Virtual (and Augmented) Reality



Ivan Sutherland: Head-mounted displays, with mechanical tracker



Oculus Rift

Scientists/Engineers need graphics too

- Visualization of various phenomena, computer aided design, virtual prototyping, simulation, etc.
- Learning at least a little bit about a graphics is highly useful for most scientists and engineers as well...

Scientific Visualization



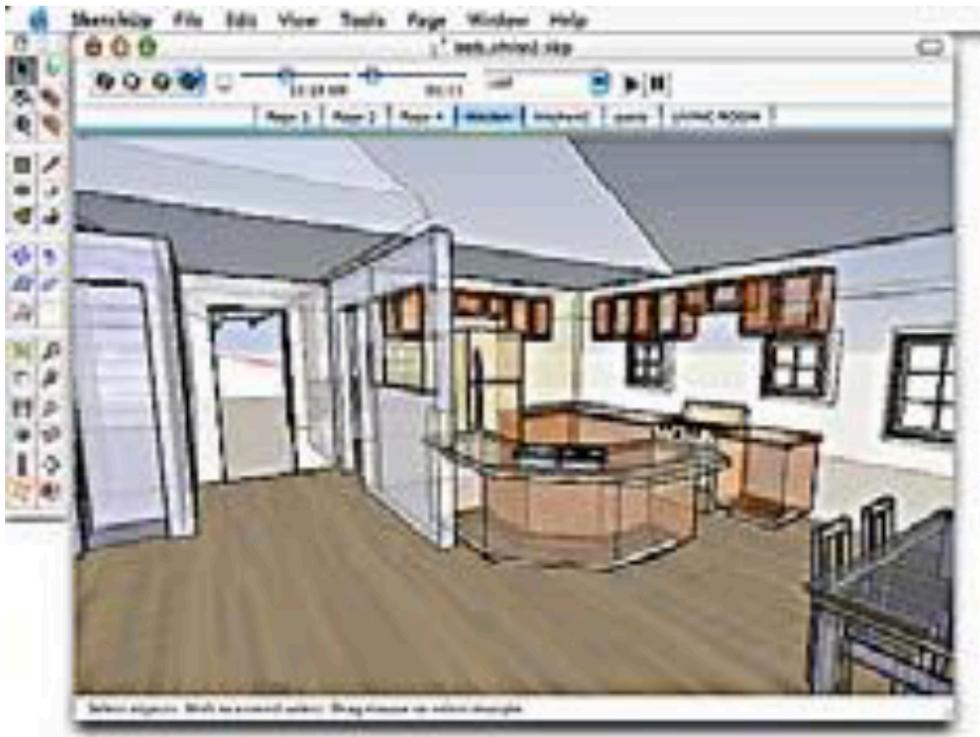
© 1995 IMDM University of Hamburg, Germany

**The Virtual Human
Karl-Heinz Hoehne**

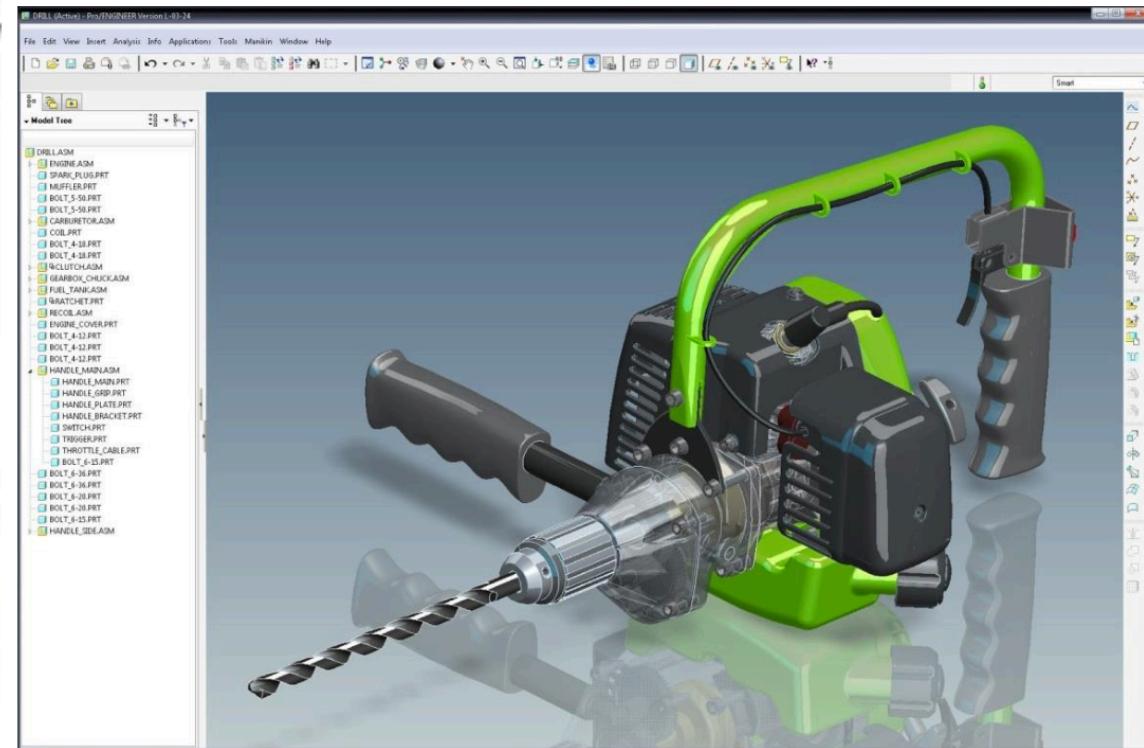


**Outside-In
The Geometry Center**

Computer-Aided Design



Sketchup



ProEngineer

Visual Simulation and Training

- Apollo spacecraft
- Flight simulators
- Driving simulators
- Surgical simulation

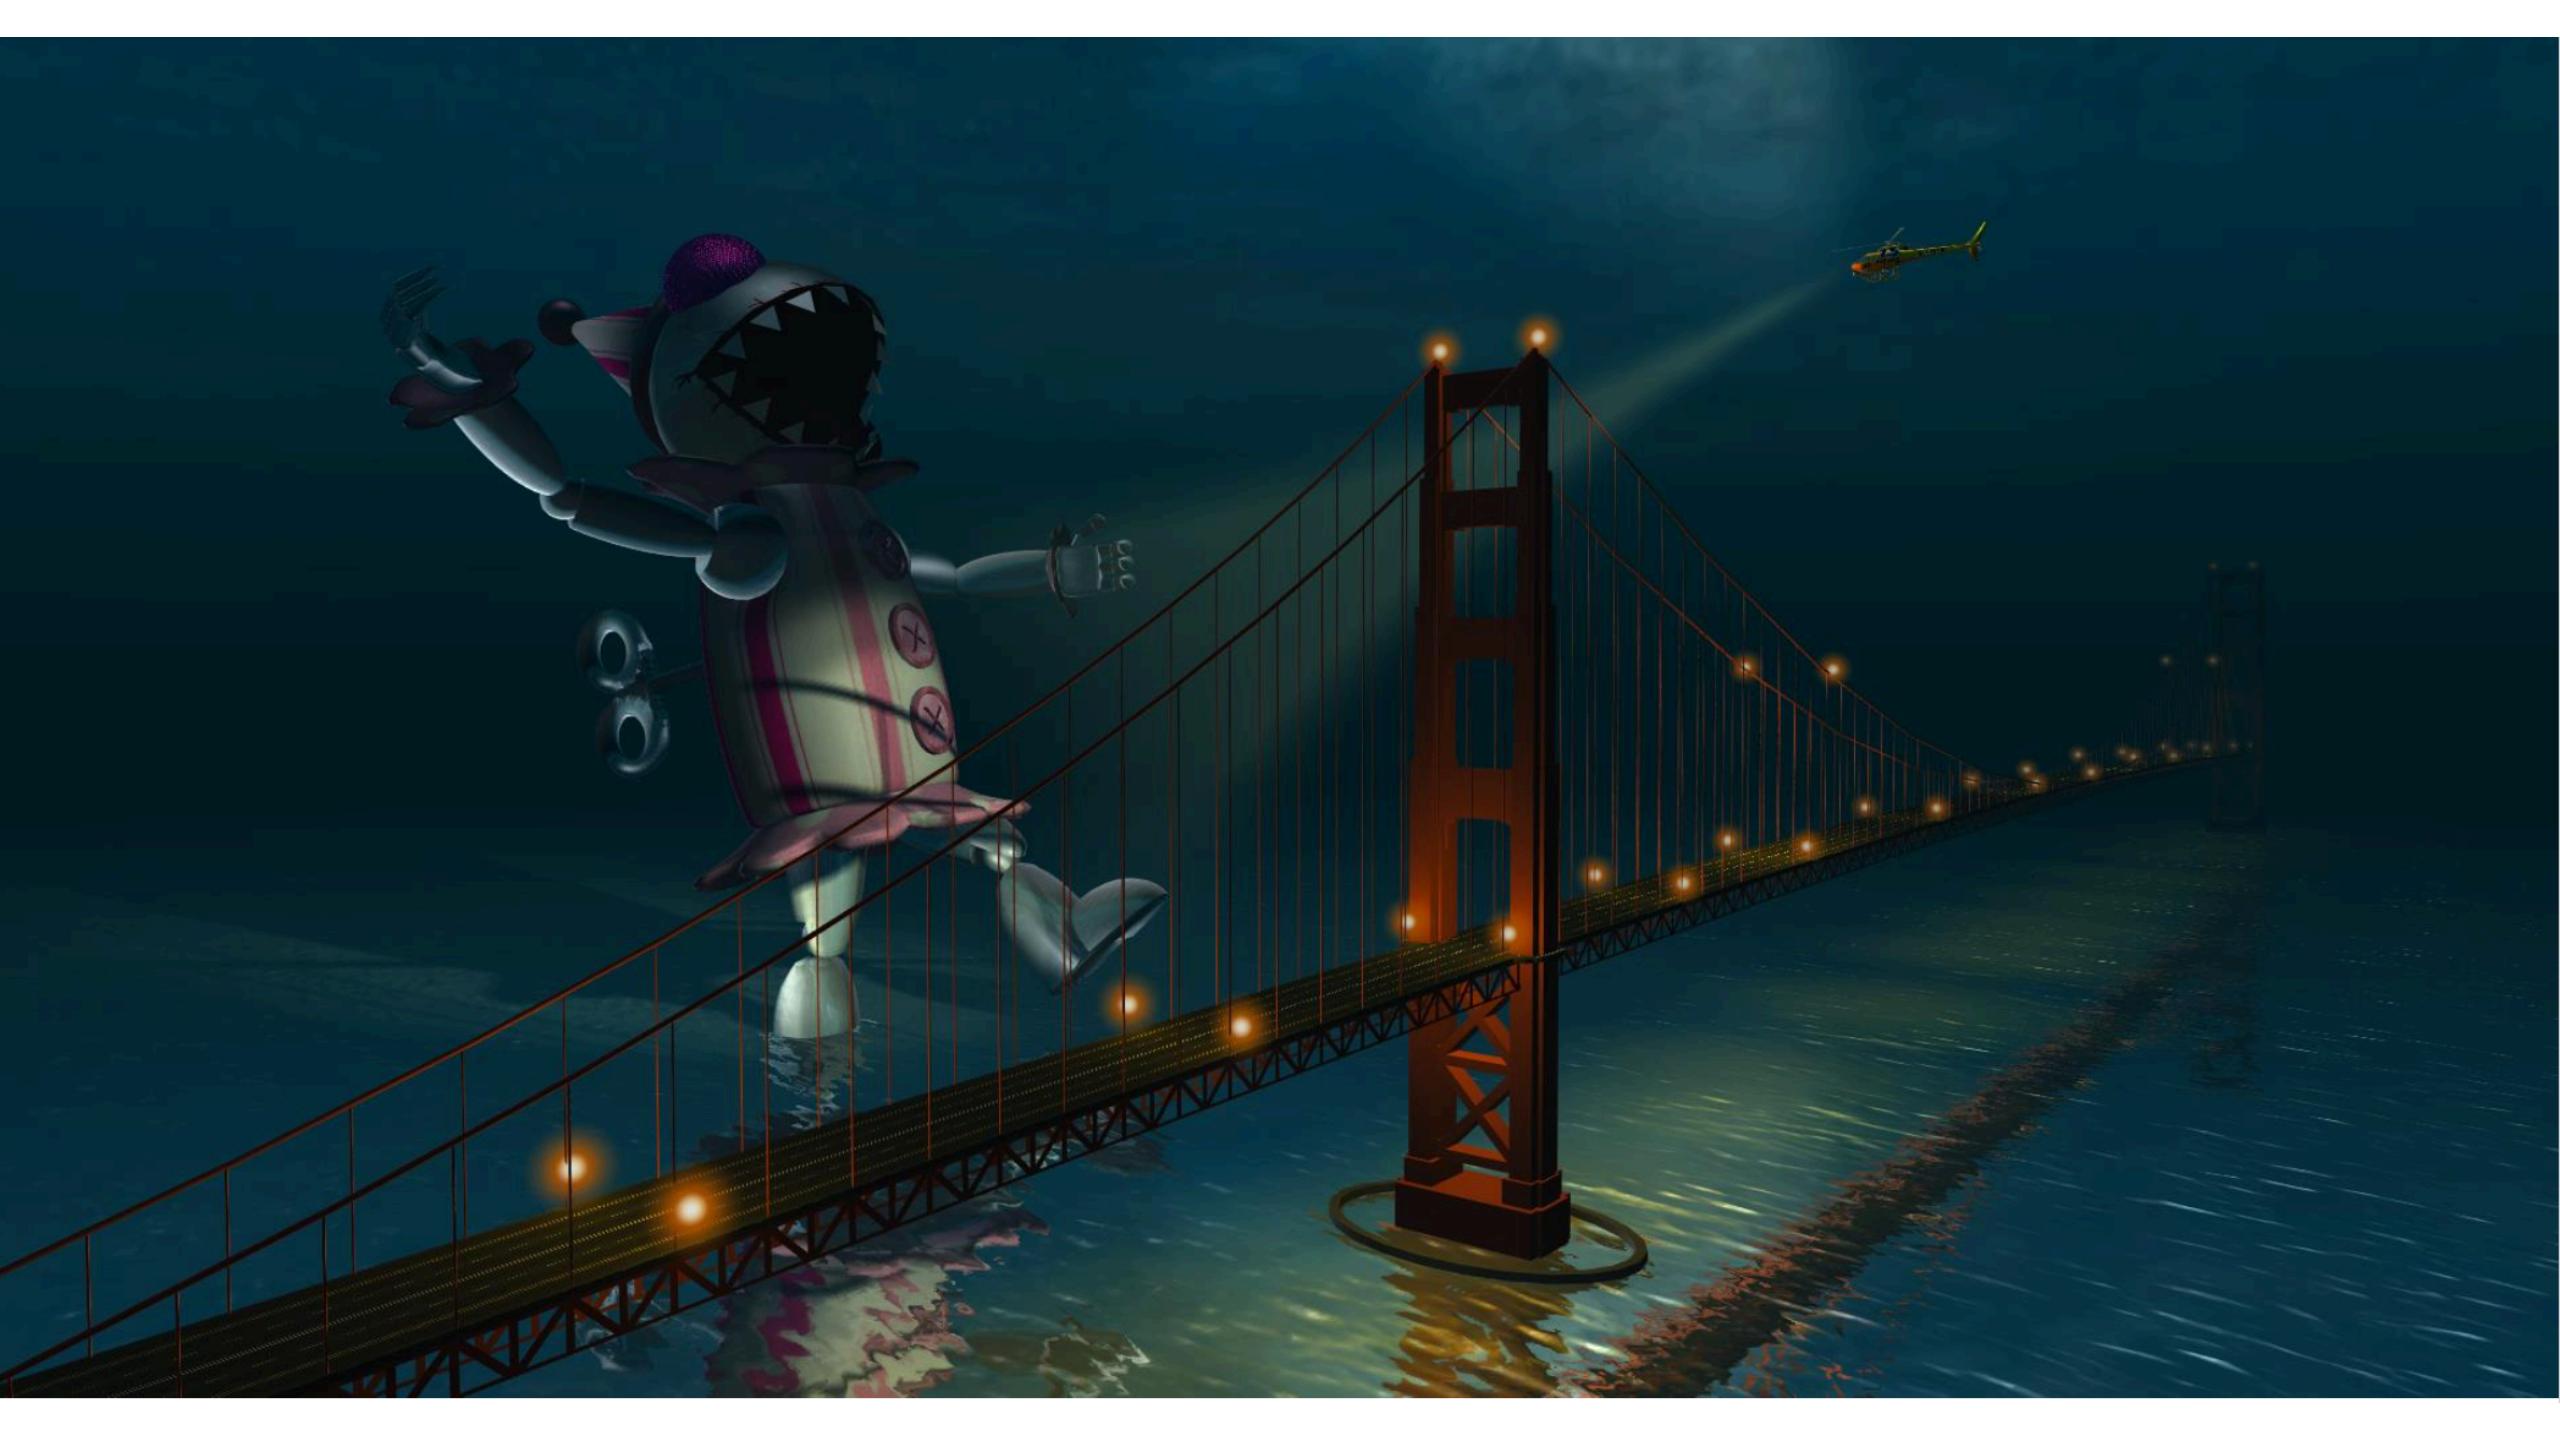


Davinci surgical robot
Intuitive Surgical

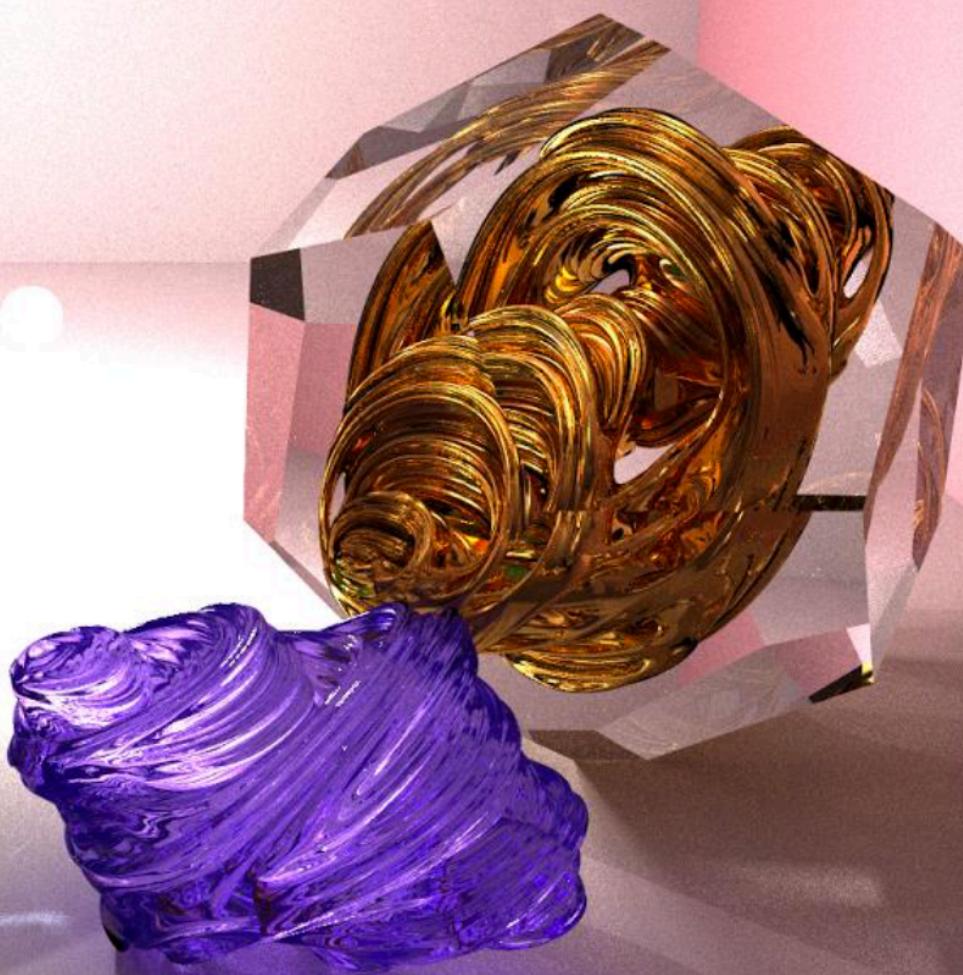
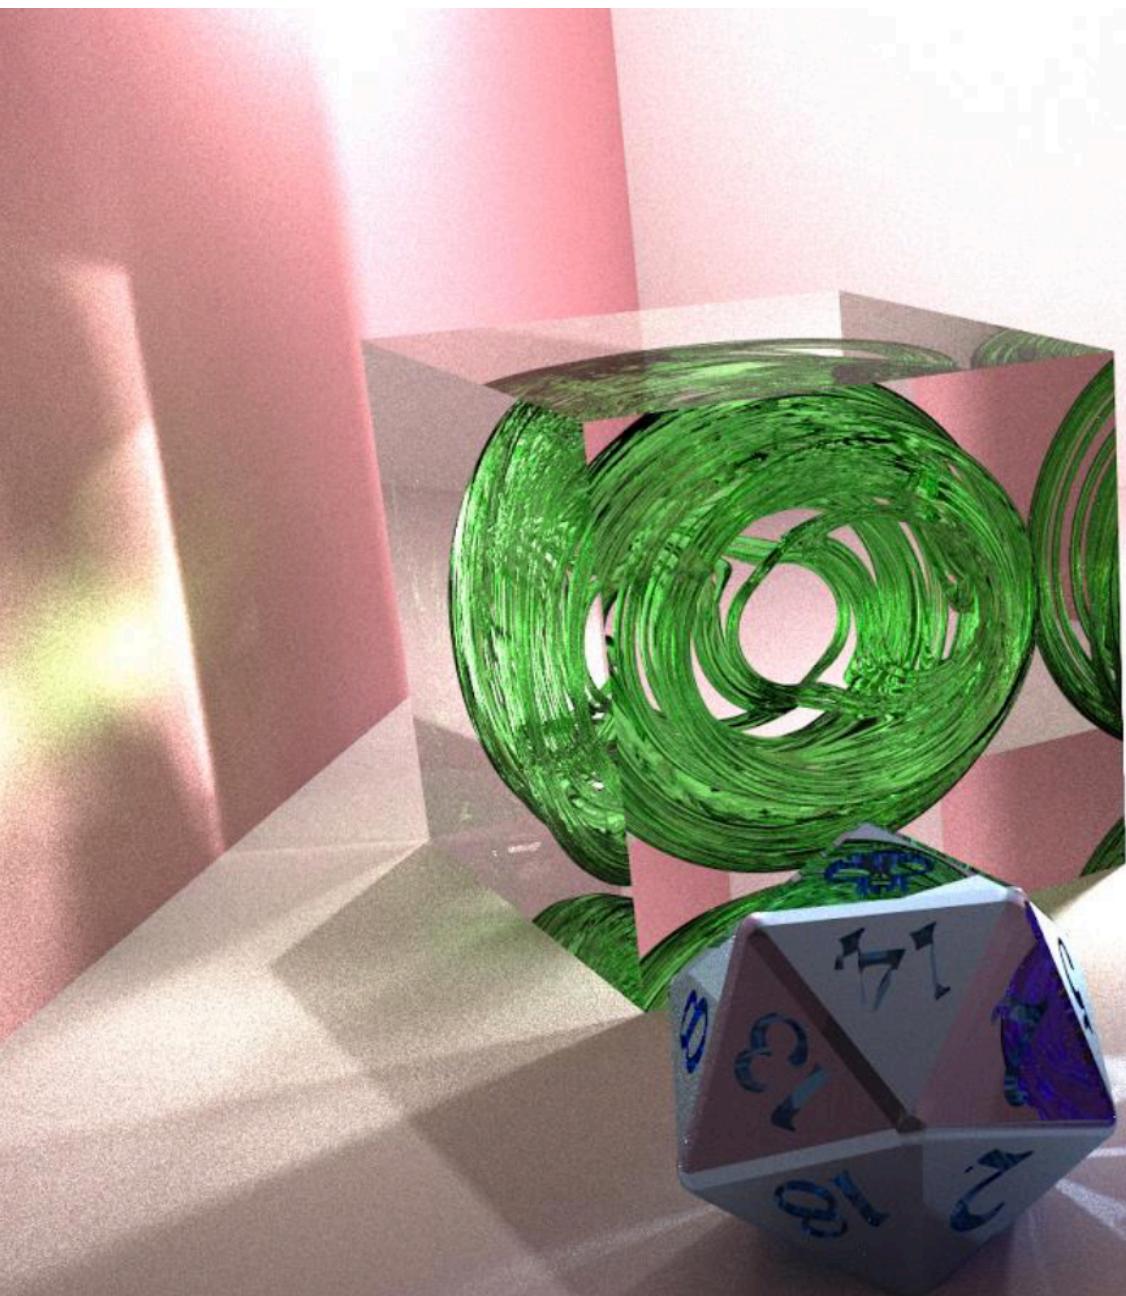


Driving simulator
Toyota Higashifuji Technical Center

More Effects









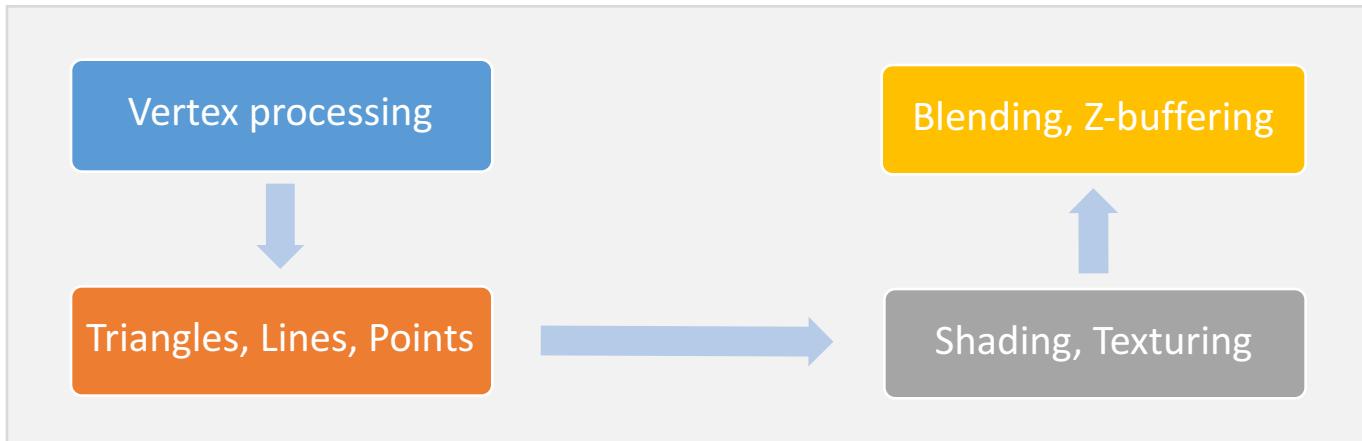
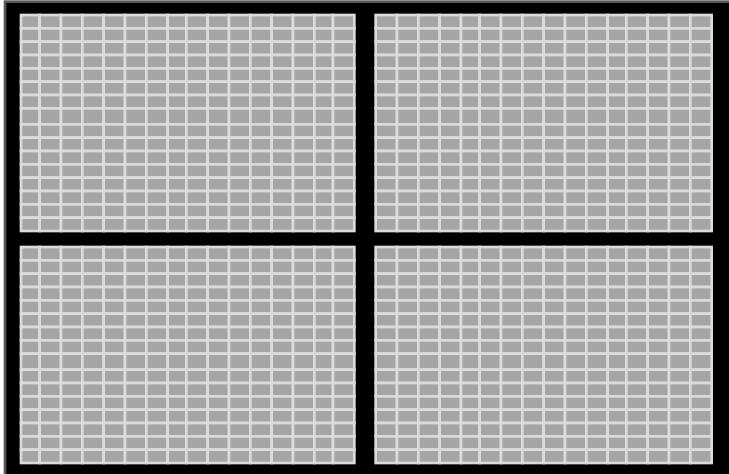






Ship in Bottle
S. S. 208

Traditional GPU workflow



GPGPU

1999-2000 computer scientists from various fields started using GPUs to accelerate a range of scientific applications.

GPU programming required the use of graphics APIs such as OpenGL and Cg.

2002 James Fung (University of Toronto) developed OpenVIDIA.

NVIDIA greatly invested in GPGPU movement and offered a number of options and libraries for a seamless experience for C, C++ and Fortran programmers.

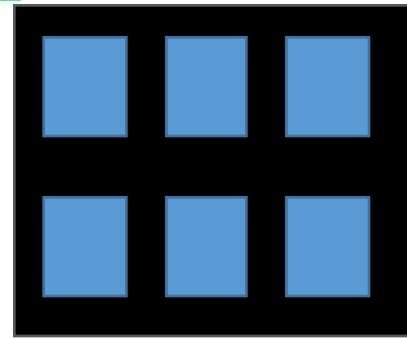
GPGPU timeline

In November 2006 Nvidia launched CUDA, an API that allows to code algorithms for execution on Geforce GPUs using C programming language.

Khronos Group defined OpenCL in 2008 supported on AMD, Nvidia and ARM platforms.

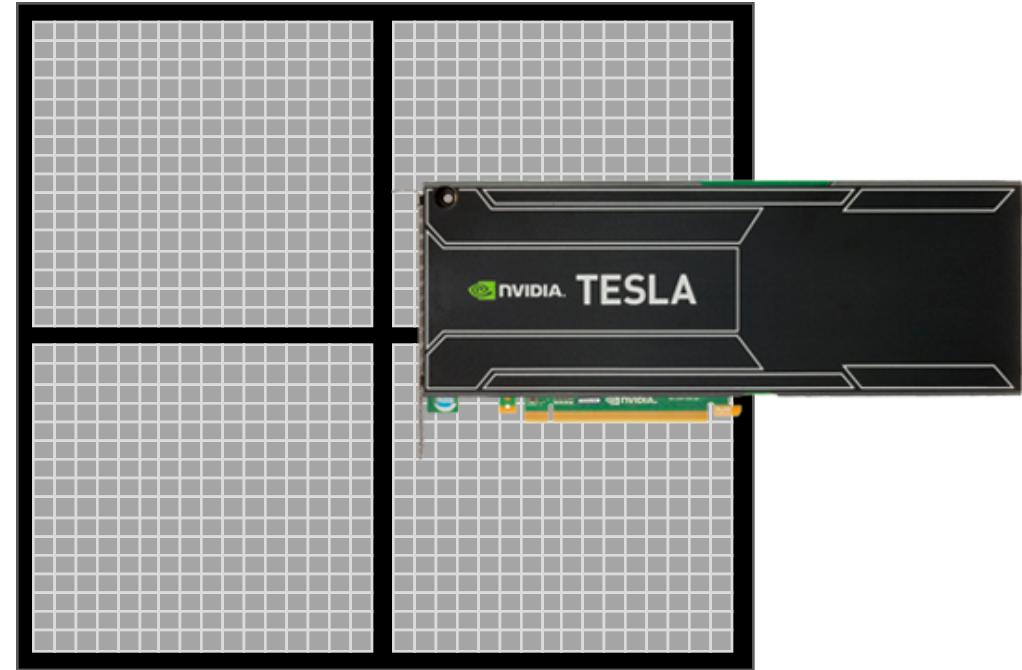
In 2012 Nvidia presented and demonstrated OpenACC - a set of directives that greatly simplify parallel programming of heterogeneous systems.

CPU



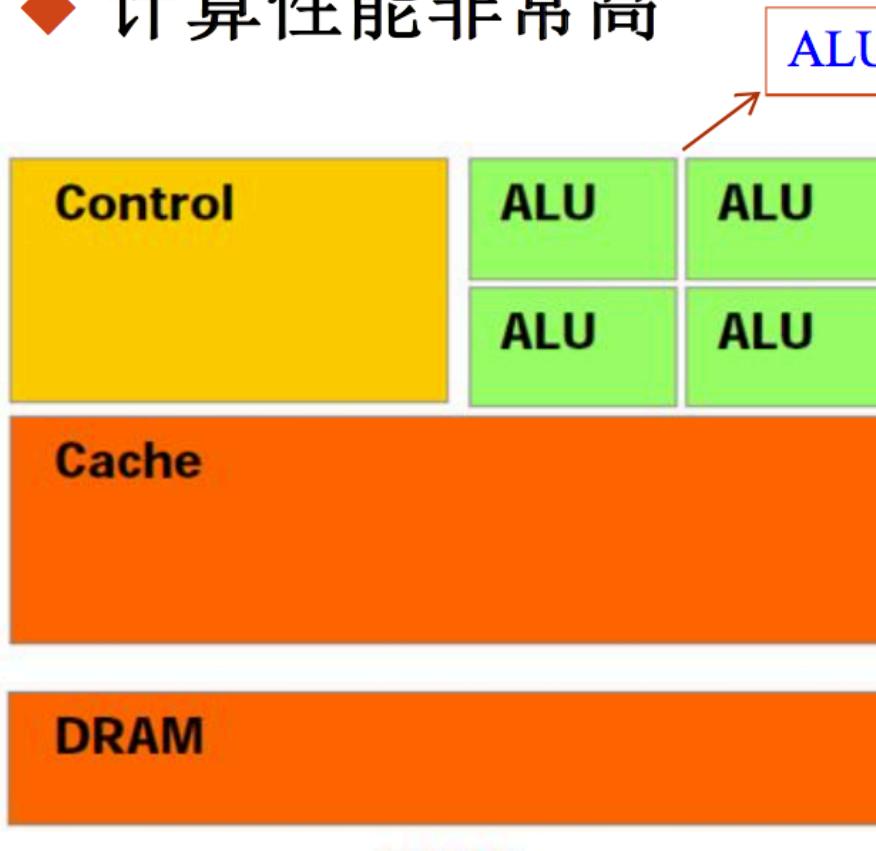
CPUs consist of a few cores optimized for serial processing

GPU



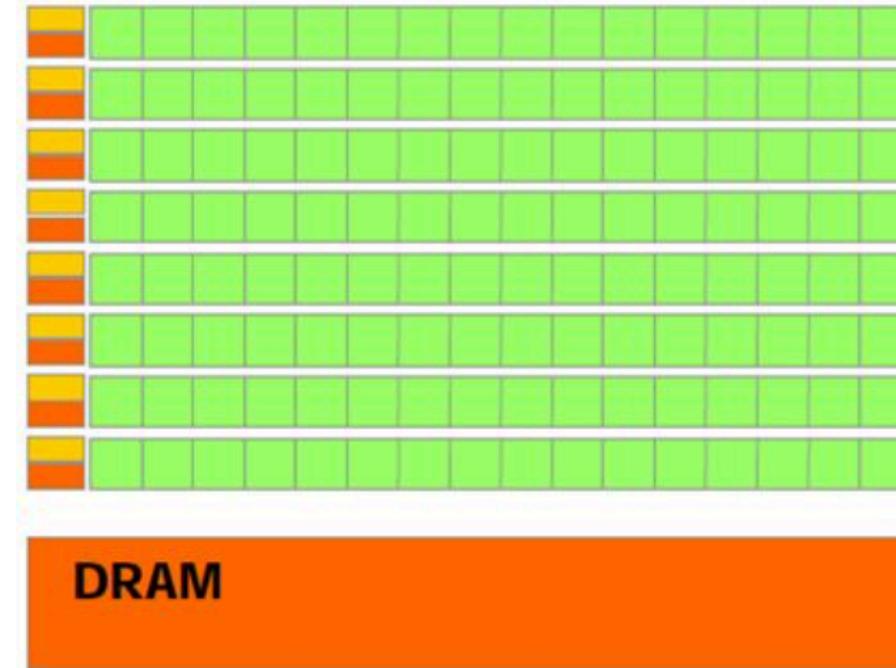
GPUs consist of hundreds or thousands of smaller, efficient cores designed for parallel performance

◆ 计算性能非常高



更多资源用于缓存

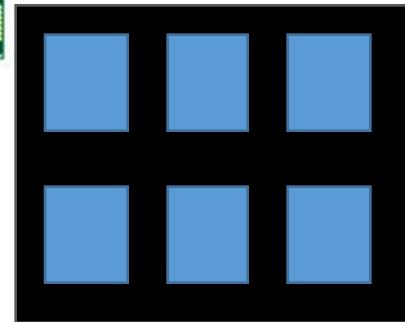
Core i7 : 4 cores 100GFlops



更多资源用于数据计算

TeslaK20 > 3.95TFlops

SCC CPU

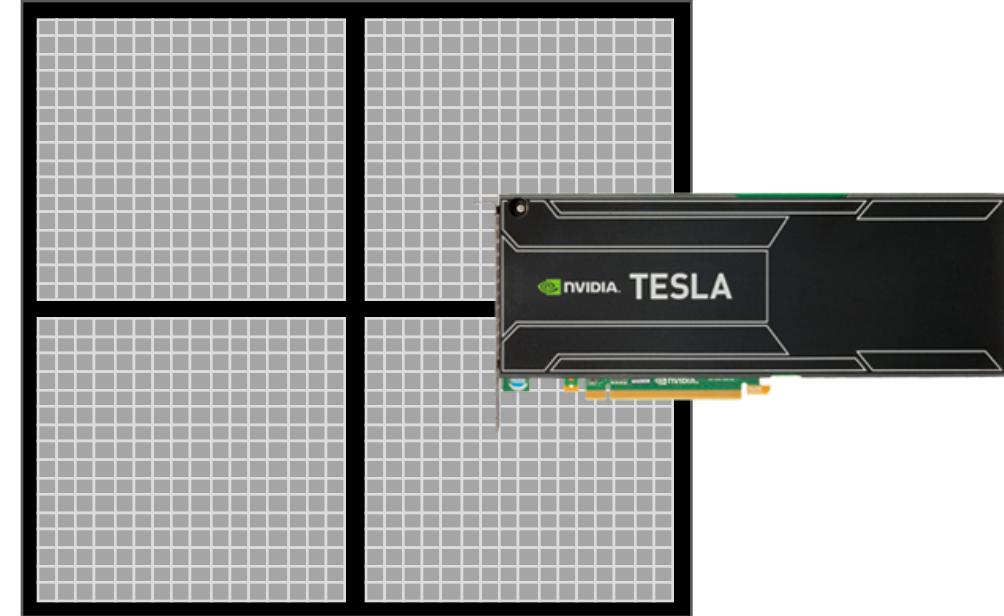


Intel Xeon X5650:

Memory size: **288 GB**

Bandwidth: **32 GB/sec**

SCC GPU

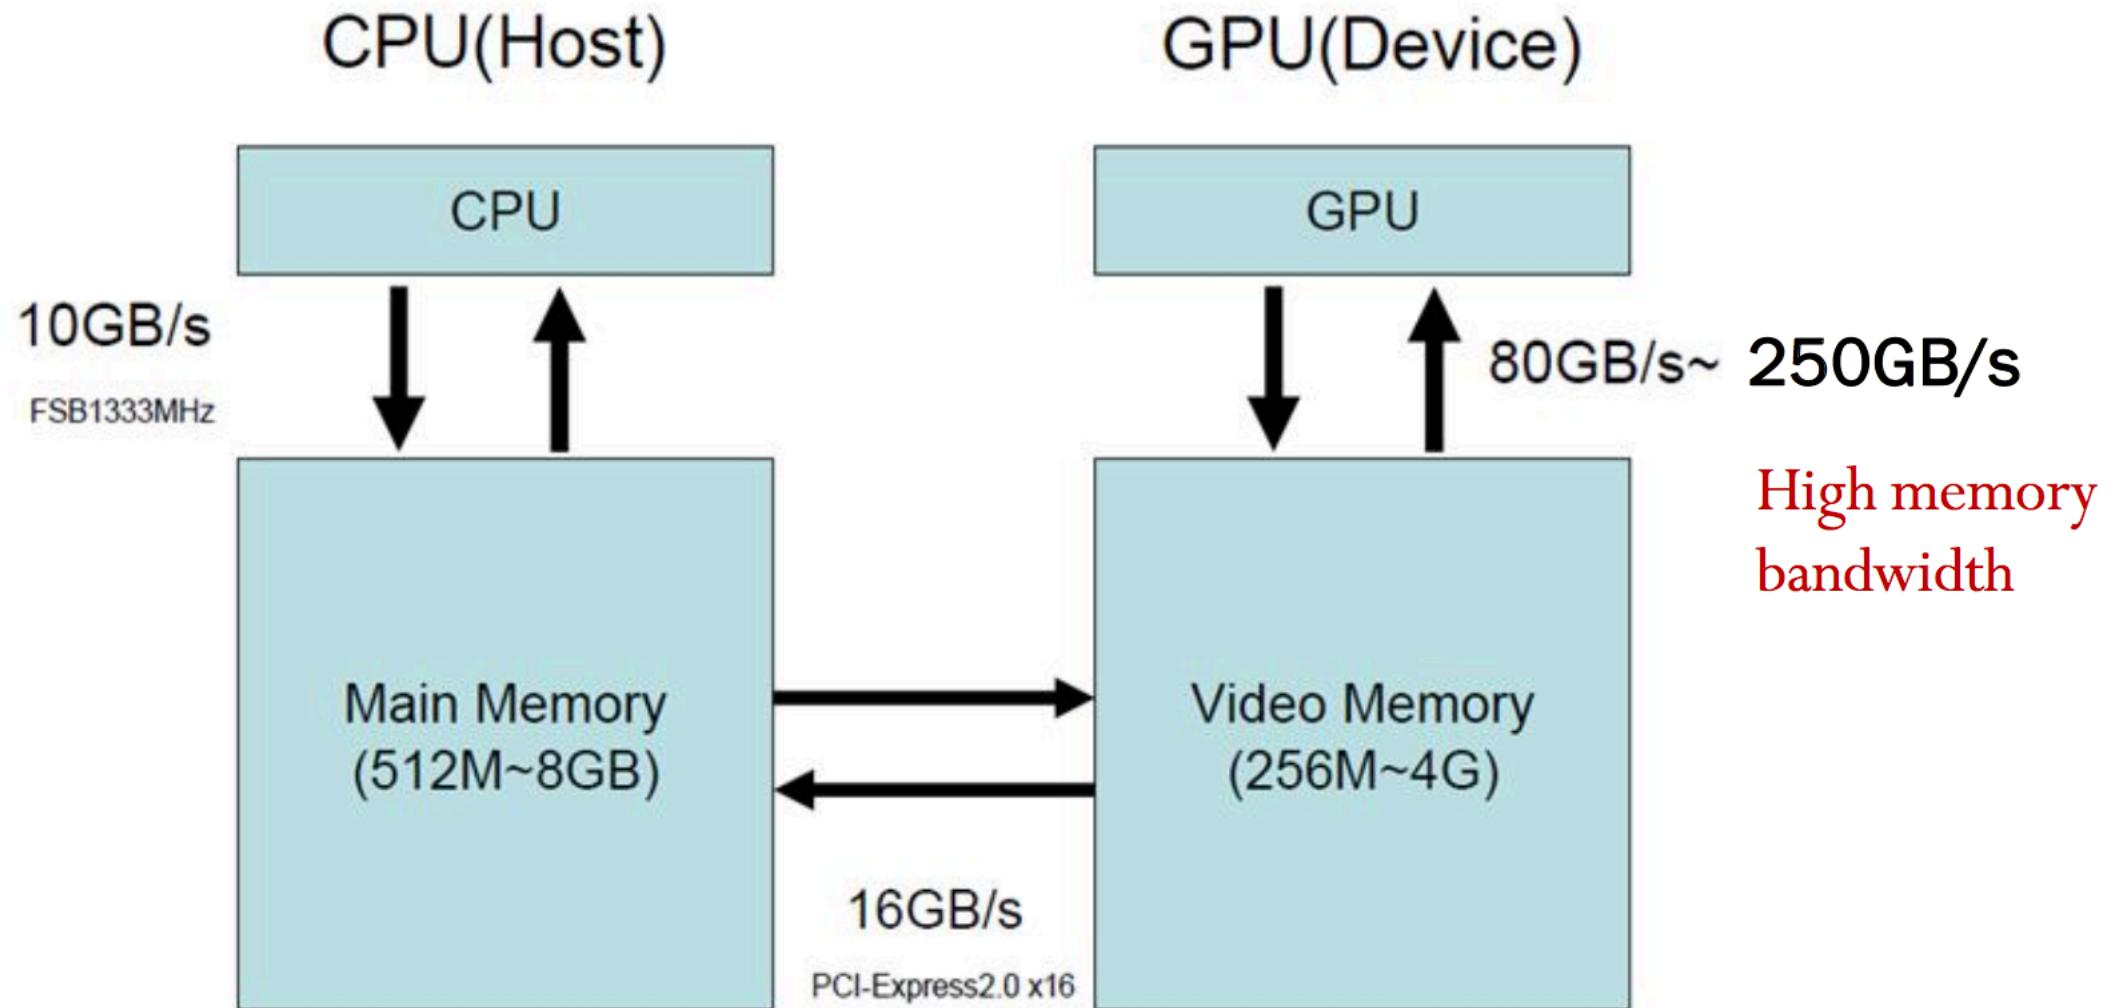


NVIDIA Tesla M2070:

Memory size: **3GB total**

Bandwidth: **150 GB/sec**

◆ 带宽非常高



GPU Computing Growth

2008

100M
CUDA-capable GPUs

150K
CUDA downloads

1
Supercomputer

4,000
Academic Papers

x 4.3

x 10.67

x 50

x 9.25

2013

430M
CUDA-capable GPUs

1.6M
CUDA downloads

50
Supercomputers

37,000
Academic Papers

GPU Computing Applications

CUDA C

OpenCL™

DirectCompute

CUDA Fortran

Application Programming Interface



NVIDIA GPU

with the CUDA Parallel Computing Architecture

App

API

GPU

MATLAB

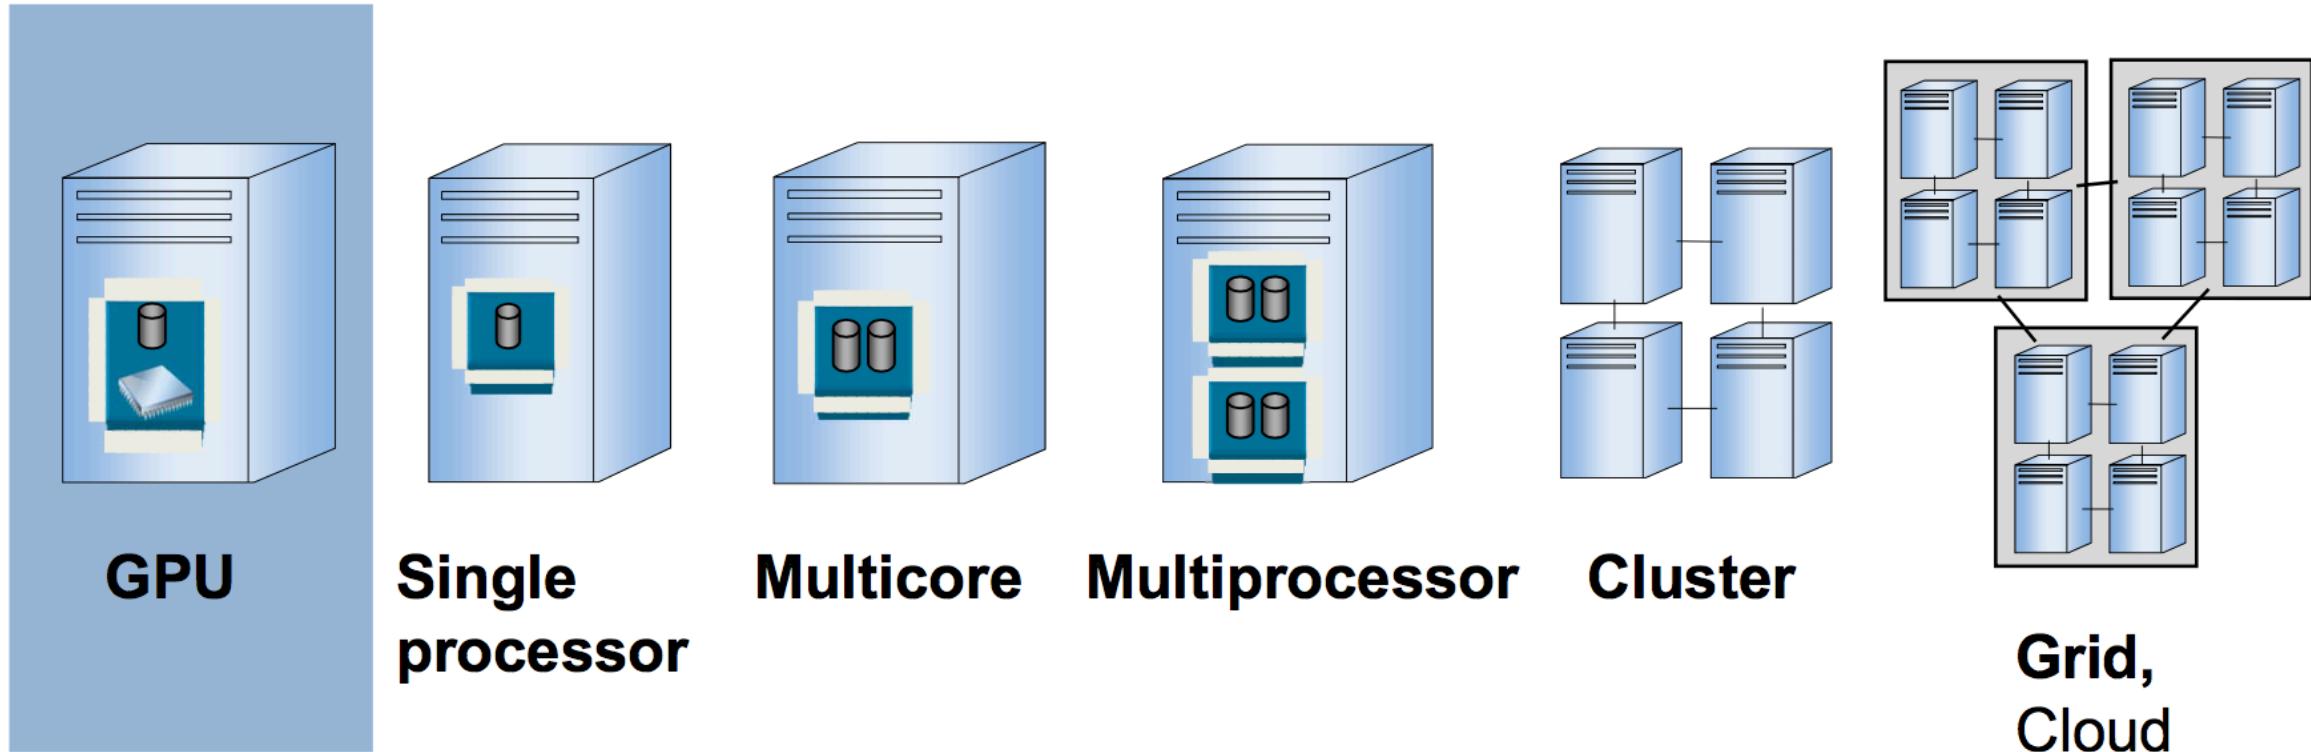
Parallel toolbox

Parallel Capabilities



Task Parallel	Data Parallel	Environment
Built-in support with Simulink, toolboxes, and blocksets		matlabpool Local workers
parfor	distributed array >200 functions	Configurations batch
job/task	spmd co-distributed array	MathWorks job manager third-party schedulers
	MPI interface	job/task

Evolving With Technology Changes



GPU Arrays

```
>> A = someArray(1000, 1000);  
>> G = gpuArray(A); % Push to GPU memory  
...  
>> F = fft(G);  
>> x = G\b;  
...  
>> z = gather(x); % Bring back into MATLAB
```

Spectrogram on the desktop (CPU only)

```
D = data;
iterations = 2000; % # of parallel iterations
stride = iterations*step; %stride of outer loop

M = ceil((numel(x)-W)/stride);%iterations needed
o = cell(M, 1); % preallocate output

for i = 1:M
    % What are the start points
    thisSP = (i-1)*stride:step: ...
        (min(numel(x)-W, i*stride)-1);

    % Move the data efficiently into a matrix
    X = copyAndWindowInput(D, window, thisSP);

    % Take lots of fft's down the colmuns
    X = abs(fft(X));

    % Return only the first part to MATLAB
    o{i} = X(1:E, 1:ratio:end);

end
```

Read my codes:

Fu_VQA_main_iteratively_W1_W2_version2_GPU.m

Spectrogram on the desktop (CPU to GPU)

```
D = data;
iterations = 2000; % # of parallel iterations
stride = iterations*step; %stride of outer loop

M = ceil((numel(x)-W)/stride);%iterations needed
o = cell(M, 1); % preallocate output

for i = 1:M
    % What are the start points
    thisSP = (i-1)*stride:step: ...
        (min(numel(x)-W, i*stride)-1);

    % Move the data efficiently into a matrix
    X = copyAndWindowInput(D, window, thisSP);

    % Take lots of fft's down the colmuns
    X = abs(fft(X));

    % Return only the first part to MATLAB
    o{i} = X(1:E, 1:ratio:end);

end
```

```
D = gpuArray(data);
iterations = 2000; % # of parallel iterations
stride = iterations*step; %stride of outer loop

M = ceil((numel(x)-W)/stride);%iterations needed
o = cell(M, 1); % preallocate output

for i = 1:M
    % What are the start points
    thisSP = (i-1)*stride:step: ...
        (min(numel(D)-W, i*stride)-1);

    % Move the data efficiently into a matrix
    X = copyAndWindowInput(D, window, thisSP);

    % Take lots of fft's down the colmuns
    X = gather(abs(fft(X)));

    % Return only the first part to MATLAB
    o{i} = X(1:E, 1:ratio:end);

end
```

Spectrogram on the desktop (GPU to parallel GPU)

```
D = gpuArray(data);
iterations = 2000; % # of parallel iterations
stride = iterations*step; %stride of outer loop

M = ceil((numel(x)-W)/stride);%iterations needed
o = cell(M, 1); % preallocate output

for i = 1:M
    % What are the start points
    thisSP = (i-1)*stride:step: ...
        (min(numel(D)-W, i*stride)-1);

    % Move the data efficiently into a matrix
    X = copyAndWindowInput(D, window, thisSP);

    % Take lots of fft's down the colmuns
    X = gather(abs(fft(X)));

    % Return only the first part to MATLAB
    o{i} = X(1:E, 1:ratio:end);

end
```

```
D = gpuArray(data);
iterations = 2000; % # of parallel iterations
stride = iterations*step; %stride of outer loop

M = ceil((numel(x)-W)/stride);%iterations needed
o = cell(M, 1); % preallocate output

parfor i = 1:M
    % What are the start points
    thisSP = (i-1)*stride:step: ...
        (min(numel(D)-W, i*stride)-1);

    % Move the data efficiently into a matrix
    X = copyAndWindowInput(D, window, thisSP);

    % Take lots of fft's down the colmuns
    X = gather(abs(fft(X)));

    % Return only the first part to MATLAB
    o{i} = X(1:E, 1:ratio:end);

end
```